

Machine Learning Project

Prediction of Wine quality Analysis

```
library(ggplot2)
library (gridExtra)

## Warning: package 'gridExtra' was built under R version 3.4.4

library(corrplot)

## Warning: package 'corrplot' was built under R version 3.4.4

## corrplot 0.84 loaded

#Load the wine quality dataset
redwine <- read.csv("C:/Users/LENOVO/Desktop/dataset/redwine.csv")
View(redwine)
redwine<-data.frame(redwine)
str(redwine)

## 'data.frame': 1599 obs. of 12 variables:
## $ fixed.acidity : num 7.4 7.8 7.8 11.2 7.4 7.4 7.9 7.3 7.8 7.5 ...
## $ volatile.acidity : num 0.7 0.88 0.76 0.28 0.7 0.66 0.6 0.65 0.58
0.5 ...
## $ citric.acid : num 0 0 0.04 0.56 0 0 0.06 0 0.02 0.36 ...
## $ residual.sugar : num 1.9 2.6 2.3 1.9 1.9 1.8 1.6 1.2 2 6.1 ...
## $ chlorides : num 0.076 0.098 0.092 0.075 0.076 0.075 0.069
0.065 0.073 0.071 ...
## $ free.sulfur.dioxide : num 11 25 15 17 11 13 15 15 9 17 ...
## $ total.sulfur.dioxide: num 34 67 54 60 34 40 59 21 18 102 ...
## $ density : num 0.998 0.997 0.997 0.998 0.998 ...
## $ pH : num 3.51 3.2 3.26 3.16 3.51 3.51 3.3 3.39 3.36
3.35 ...
## $ sulphates : num 0.56 0.68 0.65 0.58 0.56 0.56 0.46 0.47 0.57
0.8 ...
## $ alcohol : num 9.4 9.8 9.8 9.8 9.4 9.4 9.4 10 9.5 10.5 ...
## $ quality : int 5 5 5 6 5 5 5 7 7 5 ...

summary(redwine)

## fixed.acidity volatile.acidity citric.acid residual.sugar
## Min. : 4.60 Min. :0.1200 Min. :0.000 Min. : 0.900
## 1st Qu.: 7.10 1st Qu.:0.3900 1st Qu.:0.090 1st Qu.: 1.900
## Median : 7.90 Median :0.5200 Median :0.260 Median : 2.200
## Mean : 8.32 Mean :0.5278 Mean :0.271 Mean : 2.539
## 3rd Qu.: 9.20 3rd Qu.:0.6400 3rd Qu.:0.420 3rd Qu.: 2.600
## Max. :15.90 Max. :1.5800 Max. :1.000 Max. :15.500
## chlorides free.sulfur.dioxide total.sulfur.dioxide
```

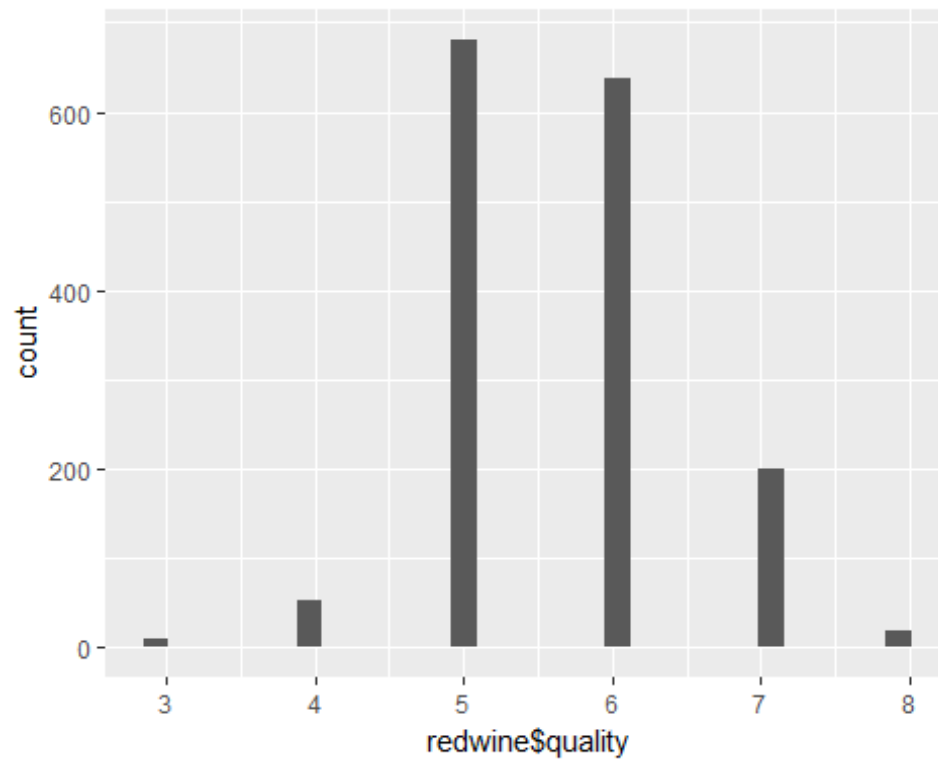
```
## Min. :0.01200 Min. : 1.00 Min. : 6.00
## 1st Qu.:0.07000 1st Qu.: 7.00 1st Qu.: 22.00
## Median :0.07900 Median :14.00 Median : 38.00
## Mean :0.08747 Mean :15.87 Mean : 46.47
## 3rd Qu.:0.09000 3rd Qu.:21.00 3rd Qu.: 62.00
## Max. :0.61100 Max. :72.00 Max. :289.00
## density pH sulphates alcohol
## Min. :0.9901 Min. :2.740 Min. :0.3300 Min. : 8.40
## 1st Qu.:0.9956 1st Qu.:3.210 1st Qu.:0.5500 1st Qu.: 9.50
## Median :0.9968 Median :3.310 Median :0.6200 Median :10.20
## Mean :0.9967 Mean :3.311 Mean :0.6581 Mean :10.42
## 3rd Qu.:0.9978 3rd Qu.:3.400 3rd Qu.:0.7300 3rd Qu.:11.10
## Max. :1.0037 Max. :4.010 Max. :2.0000 Max. :14.90
## quality
## Min. :3.000
## 1st Qu.:5.000
## Median :6.000
## Mean :5.636
## 3rd Qu.:6.000
## Max. :8.000
```

#histogram analysis

#plot shows that quality 5 is more

```
qplot(redwine$quality, geom="histogram")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
#plot shows that alcohol  
summary(redwine$alcohol)
```

```
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
##   8.40   9.50   10.20   10.42   11.10   14.90
```

```
#"alcohol" is right-skewed distributed with some outliers located at right  
side. The most frequent values are between 9.4-9.6.
```

```
qplot(redwine$alcohol, geom="histogram")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




#Now the IQR is 1.4, 3rd quantile is not very far from max value, and data gathers more in center.

```
qplot(alcohol, data=redwine, fill=I('#F79420'), binwidth =  
.1)+xlim(9.50,11.10 )
```

```
## Warning: Removed 677 rows containing non-finite values (stat_bin).
```


```
## Warning: Removed 1 rows containing missing values (geom_bar).
```



#We can see majority fixed acidity gathering in middle part

```
qplot(redwine$fixed.acidity, geom="histogram")
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



The most frequent values are between 0.4-0.6. right skewed


```
qplot(redwine$volatile.acidity, geom="histogram")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
qplot(redwine$citric.acid, geom="histogram")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
#summary
summary(redwine$citric.acid)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.000  0.090   0.260   0.271  0.420   1.000

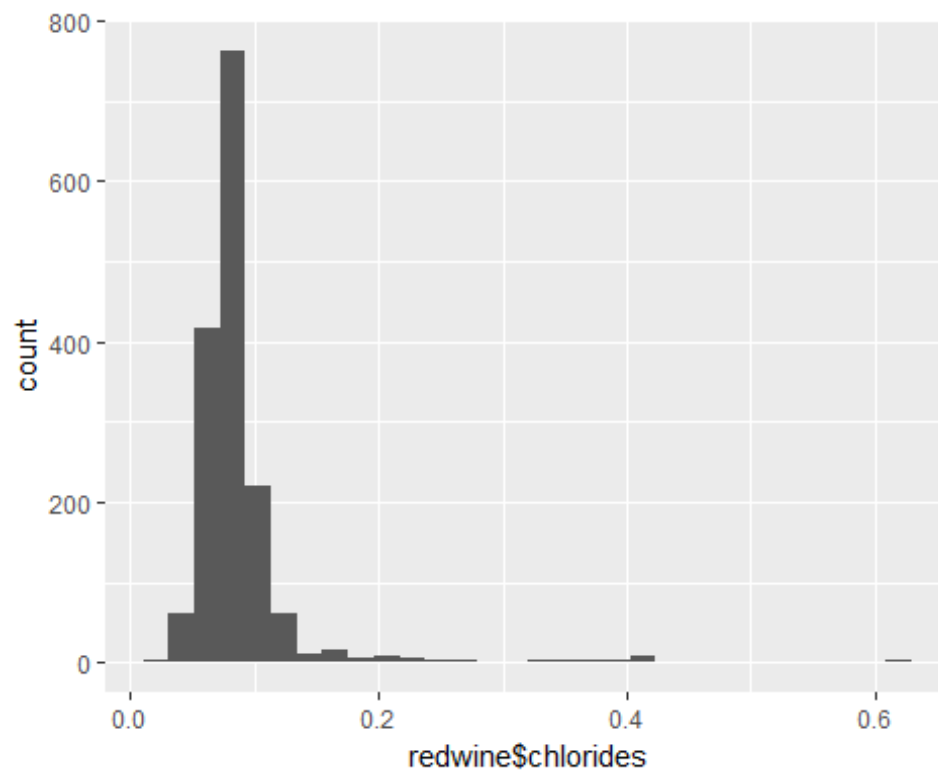
#After removing outliers. The most common citric.acid values are around 0.49
qplot(citric.acid, data=redwine, fill=I('orange'), binwidth =
.01)+xlim(0.090,.5)

## Warning: Removed 563 rows containing non-finite values (stat_bin).
```



```
qplot(redwine$chlorides, geom="histogram")


## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
qplot(chlorides, data=redwine, fill=I('orange'), binwidth = .01)+ xlim(.012, .2)
```


```
## Warning: Removed 41 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 1 rows containing missing values (geom_bar).
```




```
qplot(redwine$pH, geom="histogram")
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
qplot(pH, data=redwine, fill=I('orange'), binwidth = .01)+xlim(3.210, 3.400)
## Warning: Removed 775 rows containing non-finite values (stat_bin).
## Warning: Removed 1 rows containing missing values (geom_bar).
```



```
qplot(redwine$density, geom="histogram")
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
qplot(density, data=redwine, fill=I('orange'), binwidth = .001)
```



#The residual sugar content is high between 1.9 to 2.6. Residual sugar has a positively skewed distribution; even after eliminating the outliers distribution will remain skewed


```
qplot(redwine$residual.sugar, geom="histogram")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



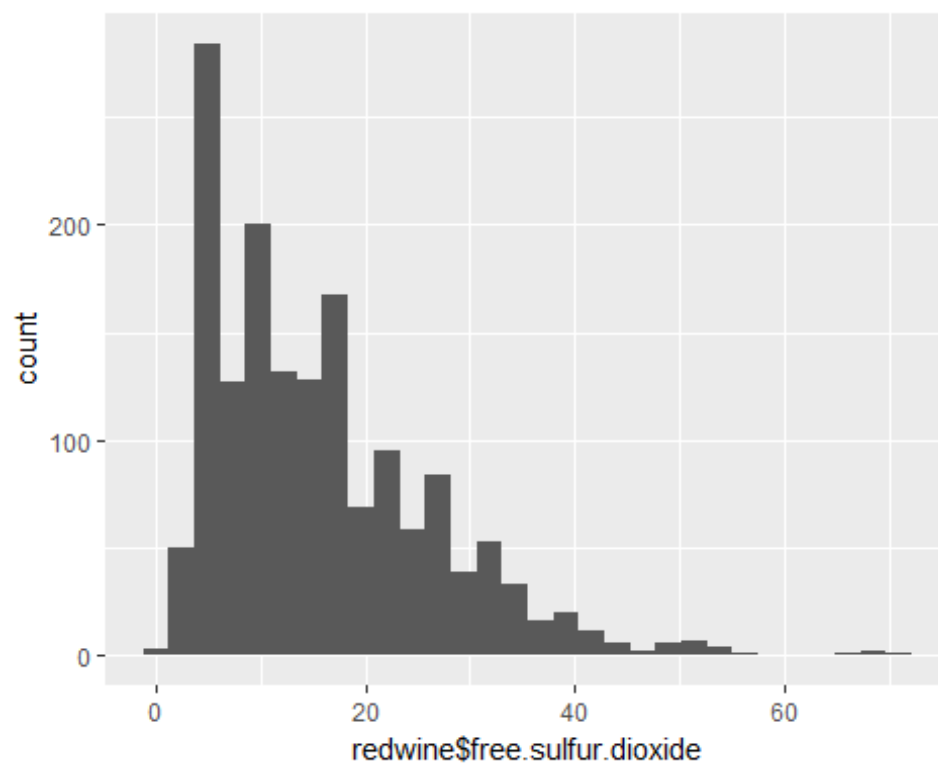
```
qplot(residual.sugar, data=redwine, fill=I('orange'), binwidth =  
.1)+xlim(1.5,5)
```

```
## Warning: Removed 134 rows containing non-finite values (stat_bin).
```




```
qplot(redwine$free.sulfur.dioxide, geom="histogram")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
qplot(redwine$total.sulfur.dioxide, geom="histogram")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
# Overall distribution
```

```
#A distribution analysis of the key paramaters pH, Total and Free  
sulfur.dioxide and Alcohol is explored, as these are the key factors
```

```
cor(redwine)
```

```
##               fixed.acidity volatile.acidity citric.acid  
## fixed.acidity      1.00000000      -0.256130895  0.67170343  
## volatile.acidity   -0.25613089      1.000000000 -0.55249568  
## citric.acid        0.67170343     -0.552495685  1.00000000  
## residual.sugar     0.11477672      0.001917882  0.14357716  
## chlorides          0.09370519      0.061297772  0.20382291  
## free.sulfur.dioxide -0.15379419     -0.010503827 -0.06097813  
## total.sulfur.dioxide -0.11318144      0.076470005  0.03553302  
## density            0.66804729      0.022026232  0.36494718  
## pH                 -0.68297819      0.234937294 -0.54190414  
## sulphates          0.18300566     -0.260986685  0.31277004  
## alcohol            -0.06166827     -0.202288027  0.10990325  
## quality            0.12405165     -0.390557780  0.22637251  
##               residual.sugar  chlorides free.sulfur.dioxide  
## fixed.acidity      0.114776724  0.093705186      -0.153794193  
## volatile.acidity    0.001917882  0.061297772      -0.010503827
```

```

## citric.acid          0.143577162  0.203822914      -0.060978129
## residual.sugar      1.000000000  0.055609535      0.187048995
## chlorides           0.055609535  1.000000000      0.005562147
## free.sulfur.dioxide 0.187048995  0.005562147      1.000000000
## total.sulfur.dioxide 0.203027882  0.047400468      0.667666450
## density             0.355283371  0.200632327     -0.021945831
## pH                  -0.085652422 -0.265026131      0.070377499
## sulphates           0.005527121  0.371260481      0.051657572
## alcohol             0.042075437 -0.221140545     -0.069408354
## quality             0.013731637 -0.128906560     -0.050656057
##
## total.sulfur.dioxide density pH
## fixed.acidity      -0.11318144  0.66804729 -0.68297819
## volatile.acidity    0.07647000  0.02202623  0.23493729
## citric.acid         0.03553302  0.36494718 -0.54190414
## residual.sugar      0.20302788  0.35528337 -0.08565242
## chlorides           0.04740047  0.20063233 -0.26502613
## free.sulfur.dioxide 0.66766645 -0.02194583  0.07037750
## total.sulfur.dioxide 1.00000000  0.07126948 -0.06649456
## density             0.07126948  1.00000000 -0.34169933
## pH                  -0.06649456 -0.34169933  1.00000000
## sulphates           0.04294684  0.14850641 -0.19664760
## alcohol             -0.20565394 -0.49617977  0.20563251
## quality             -0.18510029 -0.17491923 -0.05773139
##
## sulphates alcohol quality
## fixed.acidity      0.183005664 -0.06166827  0.12405165
## volatile.acidity   -0.260986685 -0.20228803 -0.39055778
## citric.acid        0.312770044  0.10990325  0.22637251
## residual.sugar     0.005527121  0.04207544  0.01373164
## chlorides          0.371260481 -0.22114054 -0.12890656
## free.sulfur.dioxide 0.051657572 -0.06940835 -0.05065606
## total.sulfur.dioxide 0.042946836 -0.20565394 -0.18510029
## density            0.148506412 -0.49617977 -0.17491923
## pH                 -0.196647602  0.20563251 -0.05773139
## sulphates          1.000000000  0.09359475  0.25139708
## alcohol            0.093594750  1.00000000  0.47616632
## quality            0.251397079  0.47616632  1.00000000

```

```
attach(redwine)
```

```
cor(quality,alcohol)
```

```
## [1] 0.4761663
```

```
cor(quality,pH)
```

```
## [1] -0.05773139
```

```
cor(quality,sulphates)
```


```
## [1] 0.2513971
```

```
cor(volatile.acidity,quality)
```

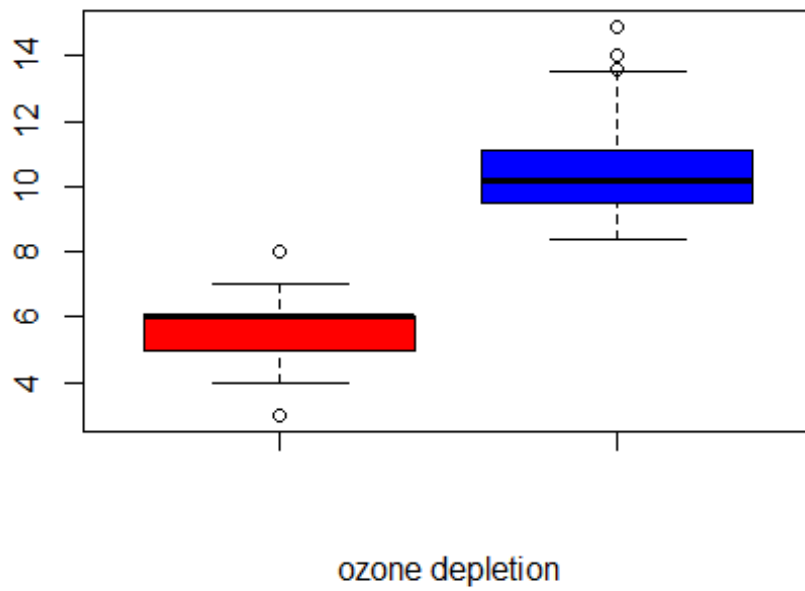


```
## [1] -0.3905578
cor(quality,citric.acid)
## [1] 0.2263725
#correlation is high for alcohol and quality


#boxplot analysis
boxplot(redwine$alcohol,main="Avg no.of rooms",col="red")
```




```
boxplot(quality,alcohol,col = c("red","blue"),xlab="ozone depletion")
```




```
ggplot(aes(factor(quality),  
            alcohol),  
       data = redwine) +  
  geom_jitter( alpha = .3) +  
  geom_boxplot( alpha = .5,color = 'blue')
```




```
ggplot(aes(factor(quality),  
            pH),  
       data = redwine) +  
  geom_jitter( alpha = .3) +  
  geom_boxplot( alpha = .5,color = 'blue')
```




```
ggplot(aes(factor(quality),  
            citric.acid),  
       data = redwine) +  
  geom_jitter( alpha = .3) +  
  geom_boxplot( alpha = .5,color = 'blue')
```




```
ggplot(aes(factor(quality),  
            volatile.acidity),  
       data = redwine) +  
  geom_jitter(alpha = .3) +  
  geom_boxplot(alpha = .5,color = 'blue')
```




```
ggplot(aes(factor(quality),  
            sulphates),  
       data = redwine) +  
  geom_jitter(alpha = .3) +  
  geom_boxplot(alpha = .5,color = 'blue')
```



```
ggplot(aes(factor(quality),  
            residual.sugar),  
       data = redwine) +  
  geom_jitter( alpha = .3) +  
  geom_boxplot( alpha = .5,color = 'blue')
```



```
ggplot(aes(factor(quality),  
            density),  
       data = redwine) +  
  geom_jitter( alpha = .3) +  
  geom_boxplot( alpha = .5,color = 'blue')
```

```
#linear regression
r<-lm(formula = quality ~alcohol, data = redwine)
r

##
## Call:
## lm(formula = quality ~ alcohol, data = redwine)
##
## Coefficients:
## (Intercept)      alcohol
##      1.8750      0.3608

summary(r)

##
## Call:
## lm(formula = quality ~ alcohol, data = redwine)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8442 -0.4112 -0.1690  0.5166  2.5888
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.87497    0.17471   10.73  <2e-16 ***
## alcohol        0.36084    0.01668   21.64  <2e-16 ***
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7104 on 1597 degrees of freedom
## Multiple R-squared:  0.2267, Adjusted R-squared:  0.2263
## F-statistic: 468.3 on 1 and 1597 DF,  p-value: < 2.2e-16
```

```
anova(r)
```

```
## Analysis of Variance Table
##
## Response: quality
##           Df Sum Sq Mean Sq F value    Pr(>F)
## alcohol     1 236.29  236.295   468.27 < 2.2e-16 ***
## Residuals 1597 805.87    0.505
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Create Training and Test data -
set.seed(100) # setting seed to reproduce results of random sampling
model <- sample(1:nrow(redwine), 0.8*nrow(redwine)) # row indices for
training data
```

```
trainingData <- redwine[model, ] # model training data
testData <- redwine[-model, ] # test data
lmMod <- lm(quality ~ alcohol, data=trainingData) # build the model
Pred <- predict(lmMod, testData)
Pred
```

```
##      3      11      13      15      20      21      23      35
## 5.400634 5.183037 5.436900 5.183037 5.183037 5.255569 5.291835 5.183037
##      37      38      43      51      57      59      64      66
## 5.763296 5.364367 5.654497 5.183037 5.328101 5.473166 5.400634 5.799562
##      70      72      82      87      89      91      92      96
## 5.654497 5.291835 5.255569 5.436900 5.291835 5.291835 5.436900 6.524885
##     100     112     114     124     133     135     138     143
## 5.110504 5.400634 5.364367 5.291835 6.561152 5.436900 5.183037 6.923813
##     144     149     151     152     157     164     173     177
## 5.255569 5.545698 5.581965 5.255569 5.654497 5.255569 5.183037 5.291835
##     187     188     194     199     215     216     219     220
## 5.183037 5.364367 5.364367 6.561152 5.400634 5.654497 5.364367 5.291835
##     222     224     233     237     238     241     244     248
## 5.255569 5.473166 5.328101 5.110504 5.183037 5.219303 5.183037 5.219303
##     253     257     258     260     266     272     275     279
## 5.654497 5.183037 5.183037 5.654497 5.545698 5.509432 5.255569 6.416087
##     281     285     287     292     293     297     311     327
## 5.219303 5.400634 5.581965 5.654497 5.581965 5.436900 5.219303 6.017159
##     329     332     335     339     343     349     361     362
## 5.364367 6.379821 6.089691 6.017159 5.400634 5.473166 5.255569 5.436900
##     364     372     374     375     381     388     395     422
## 5.545698 5.001706 5.183037 5.763296 5.255569 5.328101 5.436900 6.307288
##     428     440     443     447     448     457     458     462
## 5.690763 5.219303 5.908360 5.654497 5.908360 5.581965 5.328101 5.219303
```

##	464	468	469	474	489	503	505	513
##	5.110504	6.923813	5.183037	5.690763	5.763296	6.198490	5.763296	5.364367
##	519	531	533	534	535	543	550	558
##	6.089691	5.581965	5.581965	6.488619	5.473166	5.436900	5.255569	5.872094
##	559	563	564	571	577	589	593	594
##	6.089691	5.183037	5.509432	5.835828	5.473166	6.923813	5.364367	5.110504
##	597	600	602	608	610	611	612	617
##	5.219303	5.219303	5.110504	5.654497	6.452353	5.183037	5.255569	5.364367
##	627	630	640	647	658	661	665	667
##	5.146770	5.219303	5.581965	5.509432	5.654497	5.328101	5.545698	5.291835
##	673	674	675	687	689	694	695	703
##	5.291835	5.291835	5.219303	5.146770	5.255569	5.255569	5.255569	5.255569
##	708	710	713	717	718	727	728	738
##	5.654497	5.473166	5.255569	5.255569	5.473166	5.872094	5.291835	5.219303
##	747	749	750	756	761	762	763	768
##	5.110504	5.328101	5.291835	5.727029	5.364367	5.328101	5.473166	5.291835
##	773	777	780	782	783	786	791	796
##	5.255569	5.581965	5.400634	5.400634	5.473166	5.291835	5.364367	5.545698
##	807	811	814	821	822	824	825	826
##	6.379821	5.654497	5.980893	5.364367	6.923813	5.400634	5.581965	5.727029
##	835	838	844	853	855	856	872	878
##	5.255569	6.089691	5.255569	5.364367	5.799562	6.089691	5.690763	6.125957
##	879	880	885	891	897	899	901	911
##	5.473166	5.291835	5.473166	5.509432	6.379821	6.379821	6.125957	6.633684
##	919	924	936	940	942	946	948	952
##	5.835828	5.509432	6.125957	5.980893	6.379821	5.872094	6.343555	6.343555
##	959	964	974	976	977	986	992	1003
##	5.944626	5.908360	5.727029	5.255569	5.255569	5.944626	5.255569	6.089691
##	1007	1020	1027	1028	1030	1031	1033	1041
##	6.089691	5.618231	6.416087	6.053425	5.654497	6.017159	5.328101	5.545698
##	1043	1051	1055	1065	1068	1069	1070	1071
##	5.872094	5.618231	5.146770	6.125957	5.799562	5.799562	5.763296	5.872094
##	1074	1081	1085	1092	1094	1103	1107	1124
##	5.727029	6.125957	5.291835	5.980893	6.198490	5.908360	6.307288	5.908360
##	1143	1144	1149	1151	1154	1160	1161	1164
##	5.980893	5.473166	6.053425	6.488619	5.872094	5.763296	5.872094	5.473166
##	1175	1178	1183	1185	1204	1205	1206	1208
##	5.291835	6.452353	5.654497	5.799562	5.255569	5.835828	5.835828	5.473166
##	1218	1222	1248	1259	1268	1275	1279	1280
##	6.198490	6.017159	5.654497	5.944626	5.980893	5.908360	5.291835	6.017159
##	1281	1284	1294	1298	1299	1301	1316	1317
##	5.618231	5.545698	5.364367	6.234756	6.271022	6.379821	5.400634	6.379821
##	1319	1331	1339	1346	1347	1352	1353	1354
##	5.400634	5.400634	5.291835	5.545698	5.980893	6.125957	5.581965	5.581965
##	1355	1368	1374	1376	1378	1380	1383	1385
##	5.473166	5.654497	5.219303	5.436900	6.271022	5.763296	5.436900	5.291835
##	1386	1391	1393	1404	1407	1413	1416	1418
##	5.183037	6.379821	5.400634	5.473166	5.799562	5.799562	5.255569	6.234756
##	1439	1440	1443	1446	1448	1449	1457	1460
##	5.581965	5.860005	5.400634	5.328101	5.364367	5.364367	5.672630	6.162224

```
##      1462      1470      1472      1485      1486      1494      1499      1501
## 5.473166 5.364367 6.416087 5.799562 5.364367 5.364367 5.763296 5.328101
##      1511      1516      1524      1525      1537      1542      1544      1551
## 5.364367 5.201170 5.618231 5.473166 5.581965 5.799562 5.618231 5.291835
##      1552      1561      1562      1575      1581      1582      1590      1596
## 5.255569 5.436900 5.436900 5.654497 6.162224 5.944626 5.183037 5.908360
```

```
summary (lmMod)
```

```
##
## Call:
## lm(formula = quality ~ alcohol, data = trainingData)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8358 -0.4006 -0.1622  0.5208  2.5994
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.84655     0.19950   9.256  <2e-16 ***
## alcohol      0.36266     0.01903  19.059  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7209 on 1277 degrees of freedom
## Multiple R-squared:  0.2215, Adjusted R-squared:  0.2208
## F-statistic: 363.2 on 1 and 1277 DF,  p-value: < 2.2e-16


actuals_preds <- data.frame(cbind(actuals=testData$quality, predicted=Pred))
correlation_accuracy <- cor(actuals_preds)
correlation_accuracy

##              actuals predicted
## actuals      1.000000  0.502429
## predicted 0.502429   1.000000

head(actuals_preds)

##      actuals predicted
## 3          5  5.400634
## 11         5  5.183037
## 13         5  5.436900
## 15         5  5.183037
## 20         6  5.183037
## 21         6  5.255569

plot(testData$quality,type='l',lty=1.8,col="green")
lines(Pred,type='l',col="blue")
```



```
summary(r)
```

```
##
## Call:
## lm(formula = quality ~ alcohol, data = redwine)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8442 -0.4112 -0.1690  0.5166  2.5888
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.87497    0.17471   10.73  <2e-16 ***
## alcohol      0.36084    0.01668   21.64  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7104 on 1597 degrees of freedom
## Multiple R-squared:  0.2267, Adjusted R-squared:  0.2263
## F-statistic: 468.3 on 1 and 1597 DF,  p-value: < 2.2e-16
```

```
anova(r)
```

```
## Analysis of Variance Table
##
## Response: quality
##              Df Sum Sq Mean Sq F value    Pr(>F)
```

```

## alcohol      1 236.29 236.295 468.27 < 2.2e-16 ***
## Residuals 1597 805.87 0.505
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

#rmse value is 0.7 the model is not good so we go for classification

#knn
# converting quality into a binary factor
for (i in 1:nrow(redwine)) {
  if (redwine$quality[i] > 5)
    redwine$label[i] <- 1
  else
    redwine$label[i] <- 0
}
redwine$label <- factor(redwine$label, levels = c(0, 1), labels = c("bad",
"good"))
# removing the quality and type variable
redwine$quality <- NULL
redwine$type <- NULL

# using a subset of 1000 obs for the training set
test_indices <- sample(1:nrow(redwine), 1000)
test <- redwine[test_indices,]
train <- redwine[-test_indices,]

set.seed(10)
library(class)
# using 20 nearest neighbors
knn_pred <- knn(train = train[, -12],
               test = test[, -12],
               cl = train$label,
               k = 35, prob = TRUE)

View(redwine)
# confusion matrix
knn_conf <- table(pred = knn_pred, true = test$label)
knn_conf

##           true
## pred    bad good
## bad    228  117
## good   245  410

# Creating the ROC curve for knn
# library(dplyr) is loaded
knn_prob <- attr(knn_pred, "prob")
knn_prob <- 2 * ifelse(knn_pred == "-1", 1-knn_prob, knn_prob) - 1


# confusion matrix
knn_conf <- table(pred = knn_pred, true = test$label)

```

```

knn_conf
      true
pred  bad good
bad    0    0
good   0    0
# Plotting the KNN ROC curve
> plot(knn_roc_perf, col = 2, lwd = 3,
+      main = "ROC Curve for kNN, k = 10")
> abline(0,1)
> # Area under the knn curve
> knn_auc_perf <- performance(knn_roc_pred, measure = "auc")
> knn_AUC <- knn_auc_perf@y.values[[1]]
> knn_AUC
[1] 0.55741

```



>

```

#random forest
library(randomForest)
library(class)
rf <- randomForest(formula = label ~ .,
                    data = train,
                    mtry = 8)

```

```
print(rf)
```

Call:

```
randomForest(formula = label ~ ., data = train, mtry = 8)
```

 Type of random forest: classification

 Number of trees: 500


No. of variables tried at each split: 8

 OOB estimate of error rate: 24.87%

Confusion matrix:

	bad	good	class.error
bad	194	78	0.2867647
good	71	256	0.2171254

```
varImpPlot(rf, main = "Variable Importance Plot")
```




```
rf_pred <- predict(rf, test, type = "class")
```

```
# predicting on the test set
rf_pred2 <- predict(rf, test, type = "class")
# Building the ROC Curve
> rf_pred2 <- as.data.frame(predict(rf, test, type = 'prob'))
> rf_pred_probs2 <- rf_pred2[,2]
> rf_roc_pred2 <- prediction(rf_pred_probs2, test$label)
> rf_perf2 <- performance(rf_roc_pred2,
+                          measure = "tpr",
```



```
      x.measure = "fpr")
> # Plotting the curve
> plot(rf_perf2, col = 2, lwd = 8,
+      main = "ROC Curve for randomForest with 8 variables")
rf_AUC <- rf_perf2@y.values[[1]]
> rf_AUC
[1] 0.8519437
```



The model shows that the classification method Random forest is good when compare to all other model.

